

Introduction Supervised Learning

IMA205

Pietro Gori

Deadline: Upload the answers and the two notebooks to E-campus. Please verify the deadline on E-Campus.

Theoretical questions

OLS

We have seen that the OLS estimator is equal to $\beta^* = (\mathbf{x}^T \mathbf{x})^{-1} \mathbf{x}^T \mathbf{y}$ which can be rewritten as $\beta^* = H \mathbf{y}$. Let $\tilde{\beta} = C \mathbf{y}$ be another linear unbiased estimator of β where C is a $d \times n$ matrix, e.g. $C = H + D$ where D is a non-zero matrix .

- Demonstrate that OLS is the estimator with the smallest variance: compute $\mathbf{E}[\tilde{\beta}]$ and $\text{Var}(\tilde{\beta}) = \mathbf{E}[(\tilde{\beta} - \mathbf{E}[\tilde{\beta}])(\tilde{\beta} - \mathbf{E}[\tilde{\beta}])^T]$ and show when and why $\text{Var}(\beta^*) < \text{Var}(\tilde{\beta})$. Which assumption of OLS do we need to use ?

Ridge regression

Suppose that both \mathbf{y} and the columns of \mathbf{x} are centered (\mathbf{y}_c and \mathbf{x}_c) so that we do not need the intercept β_0 . In this case, the matrix \mathbf{x}_c has d (rather than $d+1$) columns. We can thus write the criterion for ridge regression as:

$$\beta_{ridge}^* = \arg \min_{\beta} (\mathbf{y}_c - \mathbf{x}_c \beta)^T (\mathbf{y}_c - \mathbf{x}_c \beta) + \lambda \|\beta\|_2^2 \quad (1)$$

- Show that the estimator of ridge regression is biased (that is $\mathbf{E}[\beta_{ridge}^*] \neq \beta$).
- Recall that the SVD decomposition is $\mathbf{x}_c = U D V^T$. Write down by hand the solution β_{ridge}^* using the SVD decomposition. When is it useful using this decomposition ? Hint: do you need to invert a matrix ?
- Remember that $\text{Var}(\beta_{OLS}^*) = \sigma^2 (\mathbf{x}^T \mathbf{x})^{-1}$. Show that $\text{Var}(\beta_{OLS}^*) \geq \text{Var}(\beta_{ridge}^*)$.
- When λ increases what happens to the bias and to the variance ? Hint: Compute $\text{MSE} = \mathbf{E}[(y_0 - x_0^T \beta_{ridge}^*)^2]$ at the test point (x_0, y_0) with $y_0 = x_0^T \beta + \epsilon_0$ being the true model and $x_0^T \beta_{ridge}^*$ the ridge estimate.
- Show that $\beta_{ridge}^* = \frac{\beta_{OLS}^*}{1+\lambda}$ when $\mathbf{x}_c^T \mathbf{x}_c = I_d$

Elastic Net

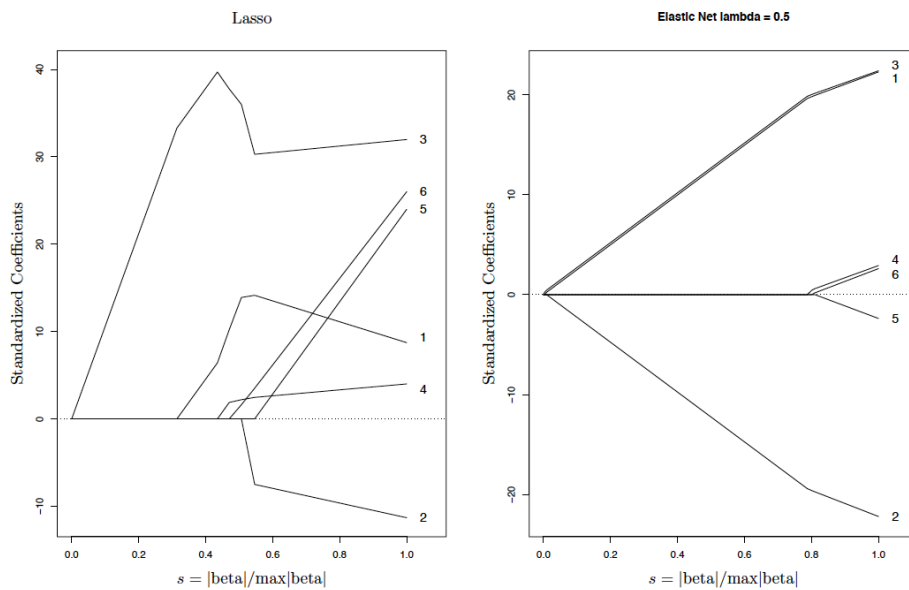
Using the previous notation, we can also combine Ridge and Lasso in the so-called Elastic Net regularization :

$$\beta_{ELNet}^* = \arg \min_{\beta} (\mathbf{y}_c - \mathbf{x}_c \beta)^T (\mathbf{y}_c - \mathbf{x}_c \beta) + \lambda_2 \|\beta\|_2^2 + \lambda_1 \|\beta\|_1 \quad (2)$$

Calling $\alpha = \frac{\lambda_2}{\lambda_1 + \lambda_2}$, solving the previous Eq. is equivalent to:

$$\beta_{ELNet}^* = \arg \min_{\beta} (\mathbf{y}_c - \mathbf{x}_c \beta)^T (\mathbf{y}_c - \mathbf{x}_c \beta) + \lambda \left[\alpha \left(\sum_{j=1}^d \beta_j^2 \right) + (1 - \alpha) \left(\sum_{j=1}^d |\beta_j| \right) \right] \quad (3)$$

- This regularization overcomes some of the limitations of the Lasso, notably:
 - If $d > N$ Lasso can select at most N variables \rightarrow ElNet removes this limitation
 - If a group of variables are highly correlated, Lasso randomly selects only one variable \rightarrow with ElNet correlated variables have a similar value (grouped)
 - Lasso solution paths tend to vary quite drastically \rightarrow ElNet regularizes the paths
 - If $N > d$ and there is high correlation between the variables, Ridge tends to have a better performance in prediction \rightarrow ElNet combines Ridge and Lasso to have better (or similar) prediction accuracy with less (or more grouped) variables



- Compute by hand the solution of Eq.2 supposing that $\mathbf{x}_c^T \mathbf{x}_c = I_d$ and show that the solution is: $\beta_{ELNet}^* = \frac{(\beta_{OLS}^*)_j \pm \frac{\lambda_1}{2}}{1 + \lambda_2}$